

Instructions.

- a. The exam is closed-book and closed-notes. You may not have your cell phone on your person.
- b. You may use algorithms done in the class as subroutines and cite their properties.
- c. Describe your algorithms completely and precisely in English, preferably using pseudo-code.
- d. Grading will be based also on the justification given and on the clarity of your arguments. Always argue correctness of your algorithms and analyze them.

Problem 1. You are given an array $A = (a_1, \dots, a_n)$ of numbers. Define μ_∞ to be the value of μ that minimizes the function

$$\max_i |a_i - \mu|$$

Design an algorithm to compute the statistic μ_∞ in time $O(n)$. You may assume that simple arithmetic and comparison operations on the a_i 's are computed each in $O(1)$ time. (10)

Solution. Let $m = \min_i a_i$ and $M = \max_i a_i$. For any value of μ ,

$$\max_i |a_i - \mu| = \max(|\mu - m|, |M - \mu|) \geq (M - m)/2 .$$

This value is attained for $\mu = (M + m)/2$. The algorithm is to find the minimum and maximum elements in the given array and return their average. This can be computed in time $O(n)$.

Problem 2. You are given a very large array $A[\cdot]$ in which the first n elements contain integers in sorted order and the rest of the cells are filled with ∞ . But you are *not* given the value of n . Describe an algorithm that takes an integer x as input and finds a position in the array containing x , if such a position exists in $O(\log n)$ time. (*Note:* You may use binary search as a routine by specifying it without writing pseudo-code.) (15)

Solution. The algorithm starts with array segment $A[1 \dots 1]$, and if x is larger than the current interval range, it keeps doubling it, that is, $[1 \dots 1], [2 \dots 3], [4 \dots 7], [8 \dots 15], \dots, [2^k, \dots, 2^{k+1} - 1]$ etc., in sequence, until it either finds the rightmost segment item $A[2^{k+1} - 1] \geq x$ or $A[2^{k+1} - 1] = \infty$. Now a binary search for x in the interval $A[2^k, \dots, 2^{k+1} - 1]$ is done. The routine $\text{BinSearch}(A, s, r, x)$ does a binary search for x in the array segment $A[s \dots r]$ and either returns *not found* or returns an index t such that $A[t] = x$ and $s \leq t \leq r$.

$\text{Search}(A, x)$

1. $k = 1, r = 1, s = 1 \ // \ r = 2^{k-1}, s = 2^k - 1$
2. **while** $A[s] < x$
3. $k = k + 1$
4. $r = 2(r + 1) - 1, s = 2s$
5. **return** $\text{BinSearch}(A, s, r, x)$.

Complexity: The outer while loop runs at most $k = \lceil \log(2N) \rceil$ times. The binary search is done over an array segment of size at most 2^k and requires $O(k)$ time, where k can be at most $\log N$. Total time is therefore $O(\log(N))$.

Problem 3. Given a binary tree on n nodes, outline an $O(n)$ time procedure that prints a leaf node that is furthest away from the root node (in terms of number of edges in a simple downward path from the root). Argue correctness of the algorithm and its time complexity. (15)

Solution.

Keep a queue Q . Enter the root node r of T into an empty queue and set $r.d = 0$. For any vertex u , $u.d$ is the shortest distance from r to u in the tree T . At each iteration until the queue is empty, dequeue a node u from the front of the queue. Enqueue all non-NIL children v of u and set $v.d = u.d + 1$. Return the last node whose dequeue makes the queue empty.

Time complexity is $O(n)$, as all edges and vertices are visited once. For correctness, note that the following invariants are maintained: $u.d$ is the length of the downward simple path from r to u . Because of the use of a queue, u 's parent is enqueued first and then u is enqueued. u is enqueued when $u.p$ is dequeued and is under processing. Hence, vertices are enqueued and dequeued respectively in order of increasing distance from the root.

Note that the d attributes are unnecessary, since, vertices are inserted as per the $.d$ values and therefore are dequeued as per the d values. Thus, the vertex dequeued at the end is one of the furthest vertices.

FURTHEST_LEAF(T)

1. $Q = \phi$
2. $r = T.root$
3. $r.d = 0$
4. $Q.ENQUEUE(r)$
5. **whilenot** ISEMPY(Q)
6. $u = Q.DEQUEUE()$
7. **if** $u.left \neq NIL$
8. $Q.ENQUEUE(u.left)$
9. $u.left.d = u.d + 1$
10. **if** $u.right \neq NIL$
11. $Q.ENQUEUE(u.right)$
12. $u.right.d = u.d + 1$
13. **return** u

Problem 4. Let A, B and C be $n \times n$ matrices. In this problem, we consider a randomized algorithm to test if $AB = C$ in expected $O(n^2)$ time.

1. Let v be an n -dimensional vector whose entries are randomly and independently chosen to be 0 or 1 with equal probability $1/2$. Prove that if M is a non-zero $n \times n$ matrix, then, $\Pr[Mv = 0] \leq 1/2$. (10)
2. Show that $\Pr[ABv = Cv] \leq 1/2$ if $AB \neq C$. Give an $O(n^2)$ expected time algorithm for checking whether $AB = C$. (10)

Solution.

1. Let M be a non-zero matrix. Then, $Mv = 0$ iff v is in the null space of M , which can be at most $n - 1$ dimensional. By construction, v is one of the 2^n end points of a unit cube anchored at the origin. So at least half of the cube vertices (or more) lie outside of any $n - 1$ dimensional plane passing through the origin. Hence, $\Pr[Mv = 0] \leq 1/2$, assuming M is non zero.
2. Note that for any $n \times n$ matrix M , and any vector u , Mu can be calculated in the standard way using $O(n^2)$ operations. Thus, $ABv = Cv$ is equivalent to $(AB - C)v = 0$ can be calculated as follows. Compute Cv in $O(n^2)$ time; compute Bv in $O(n^2)$ time and $A(Bv)$ in another $O(n^2)$ time. Thus, $ABv = Cv$ can be verified in $O(n^2)$ time.

If $AB = C$ then $ABv = Cv$ for all v . Otherwise, $\Pr[ABv = Cv] \leq 1/2$. So, we choose independently different random values of $v = v_0, v_1, \dots, v_k, \dots$. The first value for which $ABv_j \neq Cv_j$ we can conclude that $AB \neq C$.

Analysis. Suppose $AB \neq C$. The probability of getting a false positive for each of the k trials, by independence of trials, is at most $1/2^k$. So probability of a false diagnosis (i.e., $AB \neq C$ but the tests are unable to diagnose inequality) is $\leq 2^{-k}$. The expected number of trials j before we get $ABv_j \neq Cv_j$ is less or equal to $1/(1/2) = 2$.

Problem 5. Stack

Algorithm and Correctness: 16 Complexity 4

You are given a sequence S consisting of the characters left parenthesis '(' and right parenthesis ')' only and of **even** length. Find the smallest number of parenthesis reversals that are needed to make the sequence a balanced parenthesis expression. Examples

$S_1 =) ($ Output = 2, $S_2 =) () ($ Output = 2, $S_3 =) ((($ Output = 3

1. Write $S_1 =)_1 (_2$ where the subscripts are used to label the parenthesis characters. Reverse it to $(_1)_2$ to balance.
2. Let $S_2 =)_1 (_2)_3 (_4$. Minimum reversals to give balanced expression is $(_1 (_2)_3)_4$.
3. Write $S_3 =)_1 (_2 (_3 (_4$. Reverse $)_1$ to $(_1$ to give $(_1 (_2 (_3 (_4$. Then, $(_3$ and $(_4$ are each reversed to give the sequence $(_1 (_2)_3)_4$ which is balanced. Also $(_2$ and $(_4$ can be reversed to give the sequence $(_1)_2 (_3)_4$. Either way requires 3 reversals.

(Notes: The subscript numbering is shown only for convenience.)

Solution. Outline. The algorithm is as follows. Set reversals counter to 0. Scan input from left to right one character at a time.

1. If we see a $(_i$, push i onto stack.
2. If next character is a $)_j$, then, pop stack, match with whichever i was on top of stack and pair them away. However, if the stack is empty, reverse to $(_j$, push it onto stack, and increment reversals counter by one.
3. Finally, when the input is finished, the stack contains a sequence of $($ s. If this is odd, then no amount of reversals can suffice. If it is even, then change half of the stack from the top to $)$ and add this to reversals counter.

Return $reversal_count + (\text{stack height})/2$.

The correctness of the algorithm is based on the invariant that at any point of the input the minimum number of reversals is the current reversal count + $(1/2)$ stack height of unpaired left (s, if the stack height is even).

Problem 6. *Range queries over Red-Black Trees.* (20)

A *right range-query* $\text{RIGHTRANGEQ}(T, a)$ takes a red-black tree T and a key a and returns *all* the nodes x of T such that $a \leq x.key$. Note that a may or may not be present in the tree. The output order of the nodes should be as per the ordering of the inorder traversal of T . Assume T has n nodes in total and m is the number of nodes in the output of $\text{RIGHTRANGEQ}(T, a)$. Give an $O(m + \log n)$ algorithm for $\text{RIGHTRANGEQ}(T, a, b)$. (Algorithm + Correctness: 13, Complexity: 7)

Solution. Define the function $\text{CGE}(T, a)$ to be the closest greater than or equal to node x in the inorder ordering of the nodes of T such that $x.key \geq a$. In particular, if a occurs in T , then x is the earliest occurrence of key value a in the inordering. If a does not occur in T , then $x.key$ is the smallest value greater than a that occurs in the tree.

Once $x = \text{CGE}(T, a)$ is found, we can repeatedly call the SUCCESSOR function to print all nodes of T starting from x in the inordering.

How do we compute $\text{CGE}(T, a)$? First we $\text{SEARCH}(T)$ for the key a . If the answer is Yes and a node y is returned such that $y.key = a$, then, we repeatedly call PRED of y , till we obtain the first node with key $< a$. The following node gives the starting point.

If the answer to $\text{SEARCH}(T)$ for a was No, let y be the leaf node of T from where the search for a dropped off to NIL. Note that if $a > y.key$, then, our starting point is $\text{SUCC}(y)$ and otherwise if $a < y.key$, then the starting point is y itself. (Why?)

$\text{CGE}(T, a)$

1. $y = \text{NIL}, x = T$ // first search for a
2. **while** $x \neq \text{NIL}$ **and** $x.key \neq a$
3. $y = x$
4. **if** $x.key < a$ $x = x.left$
5. **else** $x = x.right$
6. **if** $x \neq \text{NIL}$ // if a is found, find the earliest a
7. $z = x; x = \text{pred}(x);$
8. **while** $x \neq \text{NIL}$ **and** $x.key == a$
9. $z = x; x = \text{pred}(x);$
10. **return** $z;$
11. **elseif** $a < y.key$
12. **return** y
13. **else return** $\text{succ}(y);$

$\text{RIGHTRANGEQ}(T, a)$

1. $x = \text{CGE}(T, a)$
2. **while** $x \neq \text{NIL}$

3. **print** x
4. $x = \text{SUCC}(x)$

Complexity Analysis: For simplicity, first assume that a is a key that appears uniquely (i.e., only once) in the tree T and the node is s . Then, `RIGHTRANGEQ` locates s along a downward path from the root and follows it by repeated calls to `SUCC`. Let S be the sequence of edges traversed from the root to s followed by the edges traversed in the repeated calls to `SUCC`.

Every edge $\{u, v\}$ is such that $u = v.p$ is traversed first in the downward direction (u, v) and some time later in the reverse direction (v, u) in S . This is because of the nature of the query: after s , all successors are printed until the next call to `succ` returns `NIL`. In this sense, every edge $\{u, v\}$ in S is traversed in both directions, downwards and later upwards.

Say that $\{u, v\}$ is traversed with $u = v.p$ bi-directionally. Then, v may or may not be printed. When is v not printed? When v is not s and is also not a successor to s . Then, v lies on the path from root to s and is a left child of its parent. The number of such v 's is at most the number of nodes on the unique path from root to s , which is bounded above by height of s . Every other edge pair (u, v) and (v, u) appears twice in S and can be charged to the vertex v that gets printed in the output. So the total number of edges traversed is at most $2h + 2m = O(\log n + m)$, since $h = O(\log n)$ for RB-trees.